

Abhiram Yenugadhati

abhiramyenugadhati@gmail.com | 669-369-9966 | [Linkedin](#) | [GitHub](#)

Experience:

Data Engineer, Amazon AGI

Mar 2025 – Present

- Built data filtering frameworks for **petabyte-scale** AGI datasets, reducing processing time by **90%** and unblocking science teams ahead of critical training deadlines
- Designed and implemented a scalable deterministic partitioning solution in **Spark** that reduced dataset processing time from 48 hours to under 6 hours, enabling 5x larger dataset handling and meeting P90 24-hour delivery SLAs
- Led data quality assessment building **AWS services** (SFNs and Lambda functions) delivering **\$4.86M** in cost savings and accelerating data acquisition timelines
- Contributed to ingestion of **160+ datasets** (4+ PB) into a S3 datalake resolving and diving deep into extraction failures, metadata issues, and processing errors at scale
- Designed and built end-to-end automation for OSHA compliance data freeze using AWS Glue, MWAA (Airflow), Redshift, and SES — replacing a 5-day manual process across 15 DAGs and multiple environments with a single orchestrated workflow

Data Engineer (Contract), Coinbase

Sep 2024 – March 2025

- Developed **monitoring** data pipelines using **Airflow** and **Snowflake** for continuous violation detection across **10+** services in real-time, reducing monitoring efforts for security DRIs (product owners and managers) by 25%
- Collaborated with **cross-functional** teams to develop dashboards using **SQL** and **Looker** for 5+ services, reducing violation detection time and improving pipeline health by 40%

Software Data Engineer (Contract), Apple

Dec 2023 – Sep 2024

- Engineered large-scale **Spark** jobs for distributed data processing, streamlining ML model training pipelines for **PyTorch**, **XGBoost** frameworks and improving data quality for downstream applications
- Developed **ELT** pipelines for People telemetry platform ingesting and monitoring **20+** data pipelines using **Splunk** and **Grafana** dashboards, improving data **reliability** and reducing processing time for downstream analytics
- Deployed **distributed** SQL query engine (Dremio) on **Kubernetes**, optimizing cross-platform queries across **S3**, **Snowflake** and increased operational efficiency in two consecutive quarters

Software Engineer- Data, Wipro Healthcare

Jun 2019 – Dec 2021

- Developed data pipelines for PHI **anonymization** leveraging **NLP** models like **SpaCy NER** for text classification and labelling, ensuring HIPAA-compliant data processing of sensitive medical records
- Developed **web-based** health data portal using Python (Flask, D3.js), Docker and Postgres maintaining FHIR standards, increasing data availability by 30%
- Developed **Jupyter notebooks** on large datasets, aggregating data from multiple sources and formats (Parquet, CSV, JSON, XML), using data analysis, mining techniques and libraries like **Pandas**, **NumPy**, **Matplotlib**, **Seaborn**, etc
- Developed **ETL** pipelines in **Pyspark** integrating **10+** data sources including **SQL**, **NoSQL** databases and **APIs**, processing semi-structured formats (Avro, JSON, XML) and ensuring robust test coverage and code quality

Projects:

Simulation of Distributed Machine Learning Platforms [PyTorch, Flask, gRPC, Websockets,] github.com/neveram/CYPHAI

Developed platform for Robo-cup soccer implementing **microservice** architecture, enabling real-time simulation (Yolov8) and communication of Sphero robots with minimal latency and enhanced gameplay experience for users

Skills:

Programming Languages: Python, SQL, Javascript

AI/ML: pandas, NumPy, sklearn, Jupyter, PyTorch, Ray, MLFlow, Langchain, RAG

Tools/ Frameworks: Spark, Kafka, Flink, Airflow, Linux, Docker, Kubernetes, Git, Splunk, Grafana, Tableau [\[DASHBOARD\]](#)

Databases & Warehouses: SQL Server, Oracle, MongoDB, Postgres, Cassandra, Snowflake, BigQuery

Certifications: Apache Airflow Champion [link](#), AWS Solution Architect [link](#)

Education:

San Jose State University – MS in Software Engineering

Jan 2022 - Dec 2023

Relevant Coursework: Enterprise Distributed Systems, Advanced Data Mining, Machine Learning

GPA: 3.7/4.0

Vardhaman College of Engineering – B.Tech in Computer Science

August 2015 - May 2019

Relevant Coursework: Data Structures, Data Mining, Operating Systems, Design Patterns, Adv Mathematics

GPA: 3.8/4.0